How Intervening on Compute Can Lead to Concrete Risk-Reducing Policies

Capacities

Visibility

- Using public information about compute quantities to estimate actors' AI capabilities.
- Required reporting of training compute usage from cloud providers and AI developers.
- International AI chip registry.
- Privacy-preserving workload monitoring.

Allocation

- Differentially advancing beneficial AI development.
- Redistributing AI development across and within countries.
- Changing the overall pace of AI progress.
- Collaborating on a joint AI megaproject.

Enforcement

- Enforcing "compute caps" by physical limits on chip-to-chip networking.
- Hardware-based remote enforcement.
- Preventing risky training runs via multiparty control.
- Digital norm enforcement.